

# Car price prediction case study

**The flow of the case study is as below:**

1. Reading the data in python
2. Defining the problem statement
3. Identifying the Target variable
4. Looking at the distribution of Target variable
5. Basic Data exploration
6. Rejecting useless columns
7. Visual Exploratory Data Analysis for data distribution (Histogram and Barcharts)
8. Feature Selection based on data distribution
9. Outlier treatment
10. Missing Values treatment
11. Visual correlation analysis
12. Statistical correlation analysis (Feature Selection)
13. Converting data to numeric for ML
14. Sampling and K-fold cross validation
15. Trying multiple Regression algorithms
16. Selecting the best Model

## **Data description**

The business meaning of each column in the data is as below

Price: The Price of the car in dollars

Age: The age of the car in months

KM: How many KMS did the car was used

FuelType: Petrol/Diesel/CNG car

HP: Horse power of the car

MetColor: Whether car has metallic color or not

Automatic: Whether car has automatic transmission or not

CC: The engine size of the car

Doors: The number of doors in the car

Weight: The weight of the car

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: data=pd.read_csv('CarPricesData.csv')
data
```

Out[3]:

	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
0	13500	23.0	46986	Diesel	90	1	0	2000.0	3	1165.0
1	13750	23.0	72937	Diesel	90	1	0	2000.0	3	1165.0
2	13950	24.0	41711	Diesel	90	1	0	2000.0	3	1165.0
3	14950	26.0	48000	Diesel	90	0	0	2000.0	3	1165.0
4	13750	30.0	38500	Diesel	90	0	0	2000.0	3	1170.0
...	...	...	...	...	...	...	...	...	...	...
1431	7500	69.0	20544	Petrol	86	1	0	1300.0	3	1025.0
1432	10845	72.0	19000	Petrol	86	0	0	1300.0	3	1015.0
1433	8500	71.0	17016	Petrol	86	0	0	1300.0	3	1015.0
1434	7250	70.0	16916	Petrol	86	1	0	1300.0	3	1015.0
1435	6950	76.0	1	Petrol	110	0	0	1600.0	5	1114.0

1436 rows × 10 columns

```
In [4]: data.shape
```

Out[4]: (1436, 10)

```
In [5]: data.head(10)
```

```
Out[5]:
```

	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
0	13500	23.0	46986	Diesel	90	1	0	2000.0	3	1165.0
1	13750	23.0	72937	Diesel	90	1	0	2000.0	3	1165.0
2	13950	24.0	41711	Diesel	90	1	0	2000.0	3	1165.0
3	14950	26.0	48000	Diesel	90	0	0	2000.0	3	1165.0
4	13750	30.0	38500	Diesel	90	0	0	2000.0	3	1170.0
5	12950	32.0	61000	Diesel	90	0	0	2000.0	3	1170.0
6	16900	27.0	94612	Diesel	90	1	0	2000.0	3	1245.0
7	18600	30.0	75889	Diesel	90	1	0	2000.0	3	1245.0
8	21500	27.0	19700	Petrol	192	0	0	1800.0	3	1185.0
9	12950	23.0	71138	Diesel	69	0	0	1900.0	3	1105.0

```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Price       1436 non-null   int64
1   Age         1434 non-null   float64
2   KM          1436 non-null   int64
3   FuelType    1432 non-null   object
4   HP          1436 non-null   int64
5   MetColor    1436 non-null   int64
6   Automatic   1436 non-null   int64
7   CC          1434 non-null   float64
8   Doors       1436 non-null   int64
9   Weight      1434 non-null   float64
dtypes: float64(3), int64(6), object(1)
memory usage: 112.3+ KB
```

```
In [7]: data.describe(include='all')
```

Out[7]:

	Price	Age	KM	FuelType	HP	MetColor	Automatic
count	1436.000000	1434.000000	1436.000000	1432	1436.000000	1436.000000	1436.000000
unique	NaN	NaN	NaN	3	NaN	NaN	NaN
top	NaN	NaN	NaN	Petrol	NaN	NaN	NaN
freq	NaN	NaN	NaN	1260	NaN	NaN	NaN
mean	10730.824513	55.986750	68533.259749	NaN	101.502089	0.674791	0.0557
std	3626.964585	18.581796	37506.448872	NaN	14.981080	0.468616	0.2294
min	4350.000000	1.000000	1.000000	NaN	69.000000	0.000000	0.0000
25%	8450.000000	44.000000	43000.000000	NaN	90.000000	0.000000	0.0000
50%	9900.000000	61.000000	63389.500000	NaN	110.000000	1.000000	0.0000
75%	11950.000000	70.000000	87020.750000	NaN	110.000000	1.000000	0.0000
max	32500.000000	80.000000	243000.000000	NaN	192.000000	1.000000	1.0000



```
In [8]: data.isna().sum() #isnull()
```

Out[8]: Price 0  
Age 2  
KM 0  
FuelType 4  
HP 0  
MetColor 0  
Automatic 0  
CC 2  
Doors 0  
Weight 2  
dtype: int64

```
In [9]: data.nunique()
```

Out[9]: Price 236  
Age 77  
KM 1263  
FuelType 3  
HP 12  
MetColor 2  
Automatic 2  
CC 12  
Doors 4  
Weight 59  
dtype: int64

```
In [10]: data['HP'].unique()
```

Out[10]: array([ 90, 192, 69, 110, 97, 71, 116, 98, 86, 72, 107, 73],  
dtype=int64)

```
In [11]: data
```

```
Out[11]:
```

	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
0	13500	23.0	46986	Diesel	90	1	0	2000.0	3	1165.0
1	13750	23.0	72937	Diesel	90	1	0	2000.0	3	1165.0
2	13950	24.0	41711	Diesel	90	1	0	2000.0	3	1165.0
3	14950	26.0	48000	Diesel	90	0	0	2000.0	3	1165.0
4	13750	30.0	38500	Diesel	90	0	0	2000.0	3	1170.0
...	...	...	...	...	...	...	...	...	...	...
1431	7500	69.0	20544	Petrol	86	1	0	1300.0	3	1025.0
1432	10845	72.0	19000	Petrol	86	0	0	1300.0	3	1015.0
1433	8500	71.0	17016	Petrol	86	0	0	1300.0	3	1015.0
1434	7250	70.0	16916	Petrol	86	1	0	1300.0	3	1015.0
1435	6950	76.0	1	Petrol	110	0	0	1600.0	5	1114.0

1436 rows × 10 columns


```
In [12]: ### EDA
```

```
In [13]: data.columns
```

```
Out[13]: Index(['Price', 'Age', 'KM', 'FuelType', 'HP', 'MetColor', 'Automatic', 'CC',  
               'Doors', 'Weight'],  
              dtype='object')
```

```
In [14]: sns.distplot(data.HP)
```

```
Out[14]: <AxesSubplot:xlabel='HP', ylabel='Density'>
```



```
In [15]: data.HP.value_counts()
```

```
Out[15]: 110      835
          86       249
          97       164
          72        73
          90        36
          69        34
          107       21
          192       11
          116        9
          98         2
          71         1
          73         1
          Name: HP, dtype: int64
```

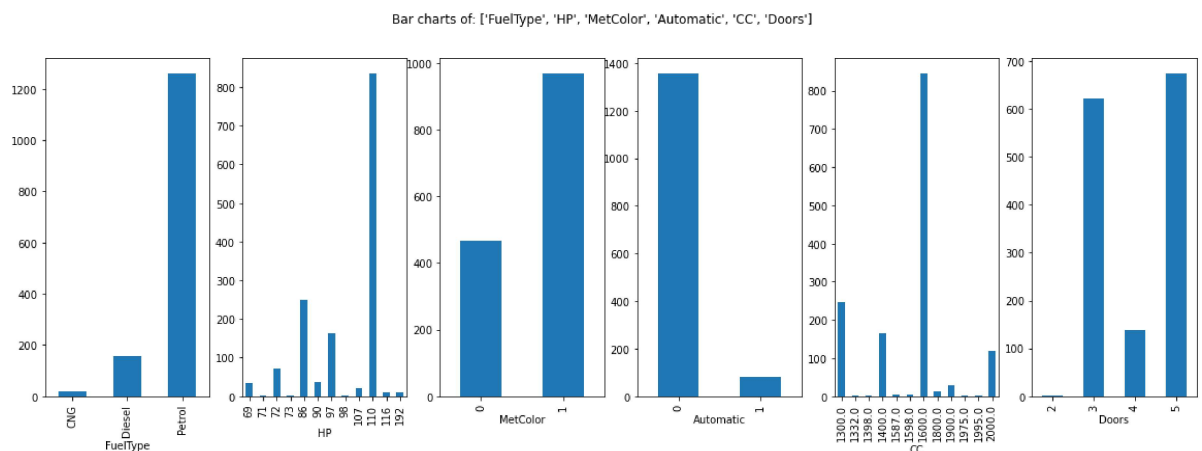
```
In [16]: import matplotlib.pyplot as plt

def plots(data, colstoplot):
    num_plots = len(colstoplot)
    fig, subplots = plt.subplots(nrows=1, ncols=num_plots, figsize=(20, 6))
    fig.suptitle('Bar charts of: ' + str(colstoplot))

    for col, subplot in zip(colstoplot, subplots):
        data.groupby(col).size().plot(kind='bar', ax=subplot)

# Example usage:
# data = ... # Provide your data here
# colstoplot = ... # Provide the columns you want to plot here
# plots(data, colstoplot)
```

```
In [17]: plots(data=data,colstoplot=['FuelType','HP','MetColor','Automatic','CC','Doors'])
```




```
In [18]: data.columns
```

```
Out[18]: Index(['Price', 'Age', 'KM', 'FuelType', 'HP', 'MetColor', 'Automatic', 'CC',
               'Doors', 'Weight'],
              dtype='object')
```


```
In [19]: sns.distplot(data.Price)
```

```
Out[19]: <AxesSubplot:xlabel='Price', ylabel='Density'>
```




```
In [20]: sns.distplot(data.Age)
```

```
Out[20]: <AxesSubplot:xlabel='Age', ylabel='Density'>
```




```
In [21]: data.Weight.hist()
```

```
Out[21]: <AxesSubplot:~>
```



```
In [22]: sns.countplot(data.CC)
```

```
Out[22]: <AxesSubplot:xlabel='CC', ylabel='count'~>
```



```
In [23]: ## outliers
```


```
data.Weight.describe()
```

```
Out[23]: count    1434.000000
mean      1072.487448
std        52.672475
min       1000.000000
25%       1040.000000
50%       1070.000000
75%       1085.000000
max       1615.000000
Name: Weight, dtype: float64
```




```
In [24]: sns.boxplot(data.Weight)
```

```
Out[24]: <AxesSubplot:xlabel='Weight'>
```




```
In [25]: data=data[data['Weight']<1150]  
sns.boxplot(data.Weight)
```

```
Out[25]: <AxesSubplot:xlabel='Weight'>
```



```
In [26]: sns.distplot(data.Weight)
```

```
Out[26]: <AxesSubplot:xlabel='Weight', ylabel='Density'>
```



```
In [27]: ##missing value treatment
```

```
In [28]: data.isna().sum()
```

```
Out[28]: Price      0
Age          2
KM           0
FuelType     3
HP           0
MetColor     0
Automatic    0
CC           2
Doors        0
Weight       0
dtype: int64
```

```
In [29]: data[data.Age.isna()]
```

```
Out[29]:
```

	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
38	15750	NaN	25329	Petrol	97	1	0	1400.0	3	1100.0
73	15750	NaN	28227	Petrol	97	1	0	1400.0	5	1110.0

```
In [30]: data.Age.fillna(0,inplace=True)
```

```
In [31]: data.isna().sum()
```

```
Out[31]: Price      0
Age      0
KM      0
FuelType  3
HP      0
MetColor  0
Automatic 0
CC      2
Doors    0
Weight   0
dtype: int64
```

```
In [32]: data.FuelType.fillna('NA',inplace=True)
```

```
In [33]: data[data.CC.isna()]
```

```
Out[33]:
```


	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
1392	7500	79.0	49827	Petrol	86	1	0	NaN	5	1035.0
1401	8950	71.0	47633	Petrol	110	1	0	NaN	5	1075.0

```
In [34]: data['CC']=np.where(data['HP']==110,1600.0,data['CC'])
data['CC']=np.where(data['HP']==86,1300.0,data['CC'])
```

```
In [ ]:
```

```
In [35]: sns.heatmap(data.corr(),annot=True)
```

```
Out[35]: <AxesSubplot:>
```



```
In [36]: data.head()
```

```
Out[36]:
```

	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
9	12950	23.0	71138	Diesel	69	0	0	1900.0	3	1105.0
17	17950	24.0	21716	Petrol	110	1	0	1600.0	3	1105.0
18	16750	24.0	25563	Petrol	110	0	0	1600.0	3	1065.0
19	16950	30.0	64359	Petrol	110	1	0	1600.0	3	1105.0
20	15950	30.0	67660	Petrol	110	1	0	1600.0	3	1105.0

```
In [37]: data=pd.get_dummies(data)
data.head()
```

```
Out[37]:
```

	Price	Age	KM	HP	MetColor	Automatic	CC	Doors	Weight	FuelType_CNG	FuelTy
9	12950	23.0	71138	69	0	0	1900.0	3	1105.0	0	
17	17950	24.0	21716	110	1	0	1600.0	3	1105.0	0	
18	16750	24.0	25563	110	0	0	1600.0	3	1065.0	0	
19	16950	30.0	64359	110	1	0	1600.0	3	1105.0	0	
20	15950	30.0	67660	110	1	0	1600.0	3	1105.0	0	



```
In [38]: x=data.drop(columns='Price')
y=data['Price']
```

```
In [39]: x.head(2)
```

```
Out[39]:
```

	Age	KM	HP	MetColor	Automatic	CC	Doors	Weight	FuelType_CNG	FuelType_Dies
9	23.0	71138	69	0	0	1900.0	3	1105.0	0	
17	24.0	21716	110	1	0	1600.0	3	1105.0	0	



```
In [40]: ## train and test split
```

```
In [41]: from sklearn.model_selection import train_test_split
```

```
In [42]: x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=
```

In [43]: x\_train

Out[43]:

	Age	KM	HP	MetColor	Automatic	CC	Doors	Weight	FuelType_CNG	FuelType_I
<b>275</b>	41.0	47350	110	1	0	1600.0	5	1075.0	0	
<b>890</b>	60.0	61100	86	1	0	1300.0	3	1015.0	0	
<b>510</b>	52.0	53561	97	1	0	1400.0	5	1060.0	0	
<b>453</b>	52.0	73000	110	0	0	1600.0	3	1055.0	0	
<b>1084</b>	73.0	134539	110	0	0	1600.0	5	1070.0	0	
...	...	...	...	...	...	...	...	...	...	...
<b>870</b>	65.0	62396	110	0	0	1600.0	3	1050.0	0	
<b>84</b>	25.0	15414	97	1	0	1400.0	5	1110.0	0	
<b>413</b>	51.0	98040	110	0	0	1600.0	5	1080.0	0	
<b>127</b>	20.0	35000	97	1	0	1400.0	5	1110.0	0	
<b>860</b>	63.0	64690	86	0	0	1300.0	3	1020.0	0	

1088 rows × 12 columns



**Thank you**