

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # Import the numpy and pandas package
```

```
import numpy as np
import pandas as pd

# Data Visualisation
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: advertising = pd.DataFrame(pd.read_csv("advertising.csv"))
advertising.head()
```

Out[3]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [4]: # 1. Missing value Treatment
```

```
advertising.isnull().sum()
```

Out[4]:

TV	0
Radio	0
Newspaper	0
Sales	0
dtype:	int64

```
In [5]: # 2. Multicollinearity
```

```
from statsmodels.stats.outliers_influence import variance_inflation_factor

x = advertising[['TV', 'Radio', 'Newspaper']]
vif_data= pd.DataFrame()
vif_data["features"] = x.columns

vif_data["VIF"] = [variance_inflation_factor(x.values,i)
                   for i in range(len(x.columns))]

print(vif_data)
```

	features	VIF
0	TV	2.486772
1	Radio	3.285462
2	Newspaper	3.055245

```
In [6]: # 3. Univariate analysis
```


```
advertising.describe()
```

Out[6]:

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

```
In [7]: sns.boxplot(advertising['Newspaper'])
```

Out[7]: <AxesSubplot:xlabel='Newspaper'>




```
In [8]: advertising = advertising[advertising['Newspaper'] < 100 ]
```

```
In [9]: advertising.shape
```

Out[9]: (198, 4)


```
In [10]: sns.boxplot(advertising['Newspaper'])
```

```
Out[10]: <AxesSubplot:xlabel='Newspaper'>
```



```
In [ ]:
```

```
In [11]: # 4.Bivariate analysis
sns.pairplot(advertising, x_vars=['TV', 'Newspaper', 'Radio'], y_vars='Sales',
plt.show())
```



```
In [12]: advertising.head()
```

```
Out[12]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [13]: # 5. Creating two data
X = advertising[['TV','Radio']]
y = advertising['Sales']
```

```
In [14]: X.head()
```

Out[14]:

	TV	Radio
0	230.1	37.8
1	44.5	39.3
2	17.2	45.9
3	151.5	41.3
4	180.8	10.8

```
In [15]: y.head()
```

Out[15]:

0	22.1
1	10.4
2	12.0
3	16.5
4	17.9

Name: Sales, dtype: float64

```
In [16]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.75, test
```

```
In [17]: X_train.shape
```

Out[17]: (148, 2)

```
In [18]: X_test.shape
```

Out[18]: (50, 2)

```
In [19]: # 6 .model building

# model 1 . Sales = F(TV)
# model 2 . sales = F(TV + Radio)
# model 3. sales = F(TV+Radio+newspaper)
```

```
In [20]: # model_1  
X_train
```

Out[20]:

	TV	Radio
187	191.1	28.7
43	206.9	8.4
68	237.4	27.5
195	38.2	3.7
76	27.5	1.6
...
108	13.1	0.4
14	204.1	32.9
93	250.9	36.5
181	218.5	5.4
104	238.2	34.3

148 rows × 2 columns

```
In [21]: import statsmodels.api as sm

# Add a constant to get an intercept
X_train_sm = sm.add_constant(X_train)

# Fit the regression line using 'OLS'
lr = sm.OLS(y_train, X_train_sm).fit()

X_train_sm
```

Out[21]:

	const	TV	Radio
187	1.0	191.1	28.7
43	1.0	206.9	8.4
68	1.0	237.4	27.5
195	1.0	38.2	3.7
76	1.0	27.5	1.6
...
108	1.0	13.1	0.4
14	1.0	204.1	32.9
93	1.0	250.9	36.5
181	1.0	218.5	5.4
104	1.0	238.2	34.3

148 rows × 3 columns

```
In [22]: print(lr.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Sales    R-squared:                  0.89
Model:                          OLS      Adj. R-squared:              0.89
Method:                        Least Squares    F-statistic:                  596.
Date:                          Tue, 11 Jul 2023    Prob (F-statistic):          1.03e-7
Time:                          22:52:09    Log-Likelihood:              -288.4
No. Observations:              148    AIC:                          582.
Df Residuals:                  145    BIC:                          591.
Df Model:                      2
Covariance Type:               nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
const                4.8447      0.349     13.870     0.000      4.154      5.535
TV                   0.0532      0.002     32.010     0.000      0.050      0.056
Radio                0.1057      0.009     11.220     0.000      0.087      0.124
=====
Omnibus:                 17.987    Durbin-Watson:              2.25
Prob(Omnibus):           0.000    Jarque-Bera (JB):           32.46
Skew:                   -0.569    Prob(JB):                   8.94e-08
Kurtosis:                4.992    Cond. No.                    43
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [23]: print(lr.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          Sales    R-squared:                0.89
Model:                  OLS      Adj. R-squared:            0.89
Method:                 Least Squares    F-statistic:          596.
Date:                   Tue, 11 Jul 2023    Prob (F-statistic):    1.03e-7
Time:                   22:52:09    Log-Likelihood:        -288.4
No. Observations:       148    AIC:                   582.
Df Residuals:           145    BIC:                   591.
Df Model:                2
Covariance Type:        nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
const                4.8447      0.349     13.870     0.000      4.154      5.535
TV                   0.0532      0.002     32.010     0.000      0.050      0.056
Radio                0.1057      0.009     11.220     0.000      0.087      0.124
=====
Omnibus:              17.987    Durbin-Watson:           2.25
Prob(Omnibus):         0.000    Jarque-Bera (JB):        32.46
Skew:                  -0.569    Prob(JB):                8.94e-08
Kurtosis:               4.992    Cond. No.                 43
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


```
In [24]: print(lr.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Sales    R-squared:                  0.89
Model:                          OLS      Adj. R-squared:              0.89
Method:                        Least Squares    F-statistic:                596.
Date:                          Tue, 11 Jul 2023    Prob (F-statistic):        1.03e-7
Time:                          22:52:09    Log-Likelihood:            -288.4
No. Observations:              148    AIC:                      582.
Df Residuals:                  145    BIC:                      591.
Df Model:                      2
Covariance Type:               nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
const                4.8447      0.349     13.870     0.000      4.154      5.535
TV                   0.0532      0.002     32.010     0.000      0.050      0.056
Radio                0.1057      0.009     11.220     0.000      0.087      0.124
=====
Omnibus:                 17.987    Durbin-Watson:              2.25
Prob(Omnibus):           0.000    Jarque-Bera (JB):           32.46
Skew:                   -0.569    Prob(JB):                   8.94e-08
Kurtosis:                4.992    Cond. No.                   43
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [25]: print(lr.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Sales    R-squared:                  0.89
Model:                          OLS      Adj. R-squared:              0.89
Method:                        Least Squares    F-statistic:                596.
Date:                          Tue, 11 Jul 2023    Prob (F-statistic):        1.03e-7
Time:                          22:52:09    Log-Likelihood:            -288.4
No. Observations:                148    AIC:                        582.
Df Residuals:                    145    BIC:                        591.
Df Model:                        2
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975
const	4.8447	0.349	13.870	0.000	4.154	5.53
TV	0.0532	0.002	32.010	0.000	0.050	0.05
Radio	0.1057	0.009	11.220	0.000	0.087	0.12

```

=====
Omnibus:                        17.987    Durbin-Watson:              2.25
Prob(Omnibus):                  0.000    Jarque-Bera (JB):           32.46
Skew:                           -0.569    Prob(JB):                   8.94e-0
Kurtosis:                       4.992    Cond. No.                    43
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [27]: y_test
```

```
Out[27]: 66      11.0
116     12.2
17      24.4
143     10.4
158      7.3
128     24.7
142     20.1
31      11.9
19      14.6
169     20.0
160     16.4
15      22.4
56       5.5
57     13.2
117      9.4
46     10.6
126      6.6
9      15.6
149     10.1
86     12.0
115     12.6
163     18.0
61     24.2
189      6.7
74     17.0
69     22.3
121      7.0
165     16.9
119      6.6
177     16.7
70     18.3
77     14.2
127     11.9
83     13.6
25     17.0
193     19.6
97     20.5
173     16.7
20     18.0
152     16.6
139     20.7
5       7.2
67     13.4
140     10.9
106      7.2
172      7.6
99     17.2
113     20.9
12      9.2
36     25.4
Name: Sales, dtype: float64
```

```
In [28]: X_test_sm = sm.add_constant(X_test)
X_test_sm
```

Out[28]:

	const	TV	Radio
66	1.0	31.5	24.6
116	1.0	139.2	14.3
17	1.0	281.4	39.6
143	1.0	104.6	5.7
158	1.0	11.7	36.9
128	1.0	220.3	49.0
142	1.0	220.5	33.2
31	1.0	112.9	17.4
19	1.0	147.3	23.9
169	1.0	284.3	10.6
160	1.0	172.5	18.1
15	1.0	195.4	47.7
56	1.0	7.3	28.1
57	1.0	136.2	19.2
117	1.0	76.4	0.8
46	1.0	89.7	9.9
126	1.0	7.8	38.9
9	1.0	199.8	2.6
149	1.0	44.7	25.8
86	1.0	76.3	27.5
115	1.0	75.1	35.0
163	1.0	163.5	36.8
61	1.0	261.3	42.7
189	1.0	18.7	12.1
74	1.0	213.4	24.6
69	1.0	216.8	43.9
121	1.0	18.8	21.7
165	1.0	234.5	3.4
119	1.0	19.4	16.0
177	1.0	170.2	7.8
70	1.0	199.1	30.6
77	1.0	120.5	28.5
127	1.0	80.2	0.0
83	1.0	68.4	44.5
25	1.0	262.9	3.5

	const	TV	Radio
193	1.0	166.8	42.0
97	1.0	184.9	21.0
173	1.0	168.4	7.1
20	1.0	218.4	27.7
152	1.0	197.6	23.3
139	1.0	184.9	43.9
5	1.0	8.7	48.9
67	1.0	139.3	14.5
140	1.0	73.4	17.0
106	1.0	25.0	11.0
172	1.0	19.6	20.1
99	1.0	135.2	41.7
113	1.0	209.6	20.6
12	1.0	23.8	35.1
36	1.0	266.9	43.8

```
In [29]: y_pred = lr.predict(X_test_sm)
```

```
In [30]: y_pred
```

```
Out[30]: 66      9.119808
116     13.755840
17      23.989330
143     11.007458
158      9.367648
128     21.735233
142     20.075518
31      12.685557
19      15.201300
169     21.077656
160     15.927673
15      20.274209
56       8.203440
57      14.114391
117      8.990432
46      10.659446
126      9.371775
9       15.740207
149      9.948333
86      11.807791
115     12.536890
163     17.426196
61      23.248616
189      7.117931
74      18.788932
69      21.010024
121      8.138141
165     17.669304
119      7.567441
177     14.716516
70      18.663106
77      14.263016
127      9.107851
83      13.185066
25      19.189514
193     18.151346
97      16.893393
173     14.546832
20      19.382440
152     17.811629
139     19.314339
5       10.476797
67      13.782300
140     10.543598
106      7.336526
172      8.011517
99      16.439893
113     18.164066
12       9.820546
36      23.662581
dtype: float64
```

```
In [31]: from sklearn.metrics import r2_score  
r_squared = r2_score(y_test, y_pred)  
r_squared
```

Out[31]: 0.9233579214065604

```
In [32]: 0.89
```

```
0.84
```

```
0.85
```

Out[32]: 0.85


```
In [33]: print(lr.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Sales    R-squared:                  0.89
Model:                          OLS      Adj. R-squared:             0.89
Method:                        Least Squares    F-statistic:                596.
Date:                          Tue, 11 Jul 2023    Prob (F-statistic):        1.03e-7
Time:                          22:52:10      Log-Likelihood:            -288.4
No. Observations:              148      AIC:                      582.
Df Residuals:                  145      BIC:                      591.
Df Model:                      2
Covariance Type:               nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
const                4.8447      0.349     13.870     0.000      4.154      5.535
TV                   0.0532      0.002     32.010     0.000      0.050      0.056
Radio                0.1057      0.009     11.220     0.000      0.087      0.124
=====
Omnibus:                 17.987    Durbin-Watson:           2.25
Prob(Omnibus):           0.000    Jarque-Bera (JB):        32.46
Skew:                   -0.569    Prob(JB):                8.94e-08
Kurtosis:                4.992    Cond. No.                43
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [34]: print(lr.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          Sales    R-squared:                0.89
Model:                  OLS      Adj. R-squared:            0.89
Method:                 Least Squares    F-statistic:          596.
Date:                   Tue, 11 Jul 2023    Prob (F-statistic):    1.03e-7
Time:                   22:52:10    Log-Likelihood:        -288.4
No. Observations:       148    AIC:                   582.
Df Residuals:           145    BIC:                   591.
Df Model:                2
Covariance Type:        nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
const                4.8447      0.349     13.870     0.000      4.154      5.535
TV                   0.0532      0.002     32.010     0.000      0.050      0.056
Radio                0.1057      0.009     11.220     0.000      0.087      0.124
=====
Omnibus:               17.987    Durbin-Watson:           2.25
Prob(Omnibus):         0.000    Jarque-Bera (JB):        32.46
Skew:                  -0.569    Prob(JB):                8.94e-08
Kurtosis:               4.992    Cond. No.                 43
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [35]: # Add a constant to X_test
X_test_sm = sm.add_constant(X_test)

# Predict the y values corresponding to X_test_sm
y_pred = lr.predict(X_test_sm)

from sklearn.metrics import r2_score
r_squared = r2_score(y_test, y_pred)
r_squared
```

Out[35]: 0.9233579214065604

```
In [36]: X_test.head()
```

Out[36]:

	TV	Radio
66	31.5	24.6
116	139.2	14.3
17	281.4	39.6
143	104.6	5.7
158	11.7	36.9

```
In [37]: y_test.head()
```

Out[37]:

66	11.0
116	12.2
17	24.4
143	10.4
158	7.3

Name: Sales, dtype: float64

```
In [38]: X_test_sm = sm.add_constant(X_test)

# Predict the y values corresponding to X_test_sm
y_pred = lr.predict(X_test_sm)

y_pred.head()
```

Out[38]:

66	9.119808
116	13.755840
17	23.989330
143	11.007458
158	9.367648

dtype: float64

```
In [39]: from sklearn.metrics import r2_score
r_squared = r2_score(y_test, y_pred)
r_squared
```

Out[39]: 0.9233579214065604

Thankyou